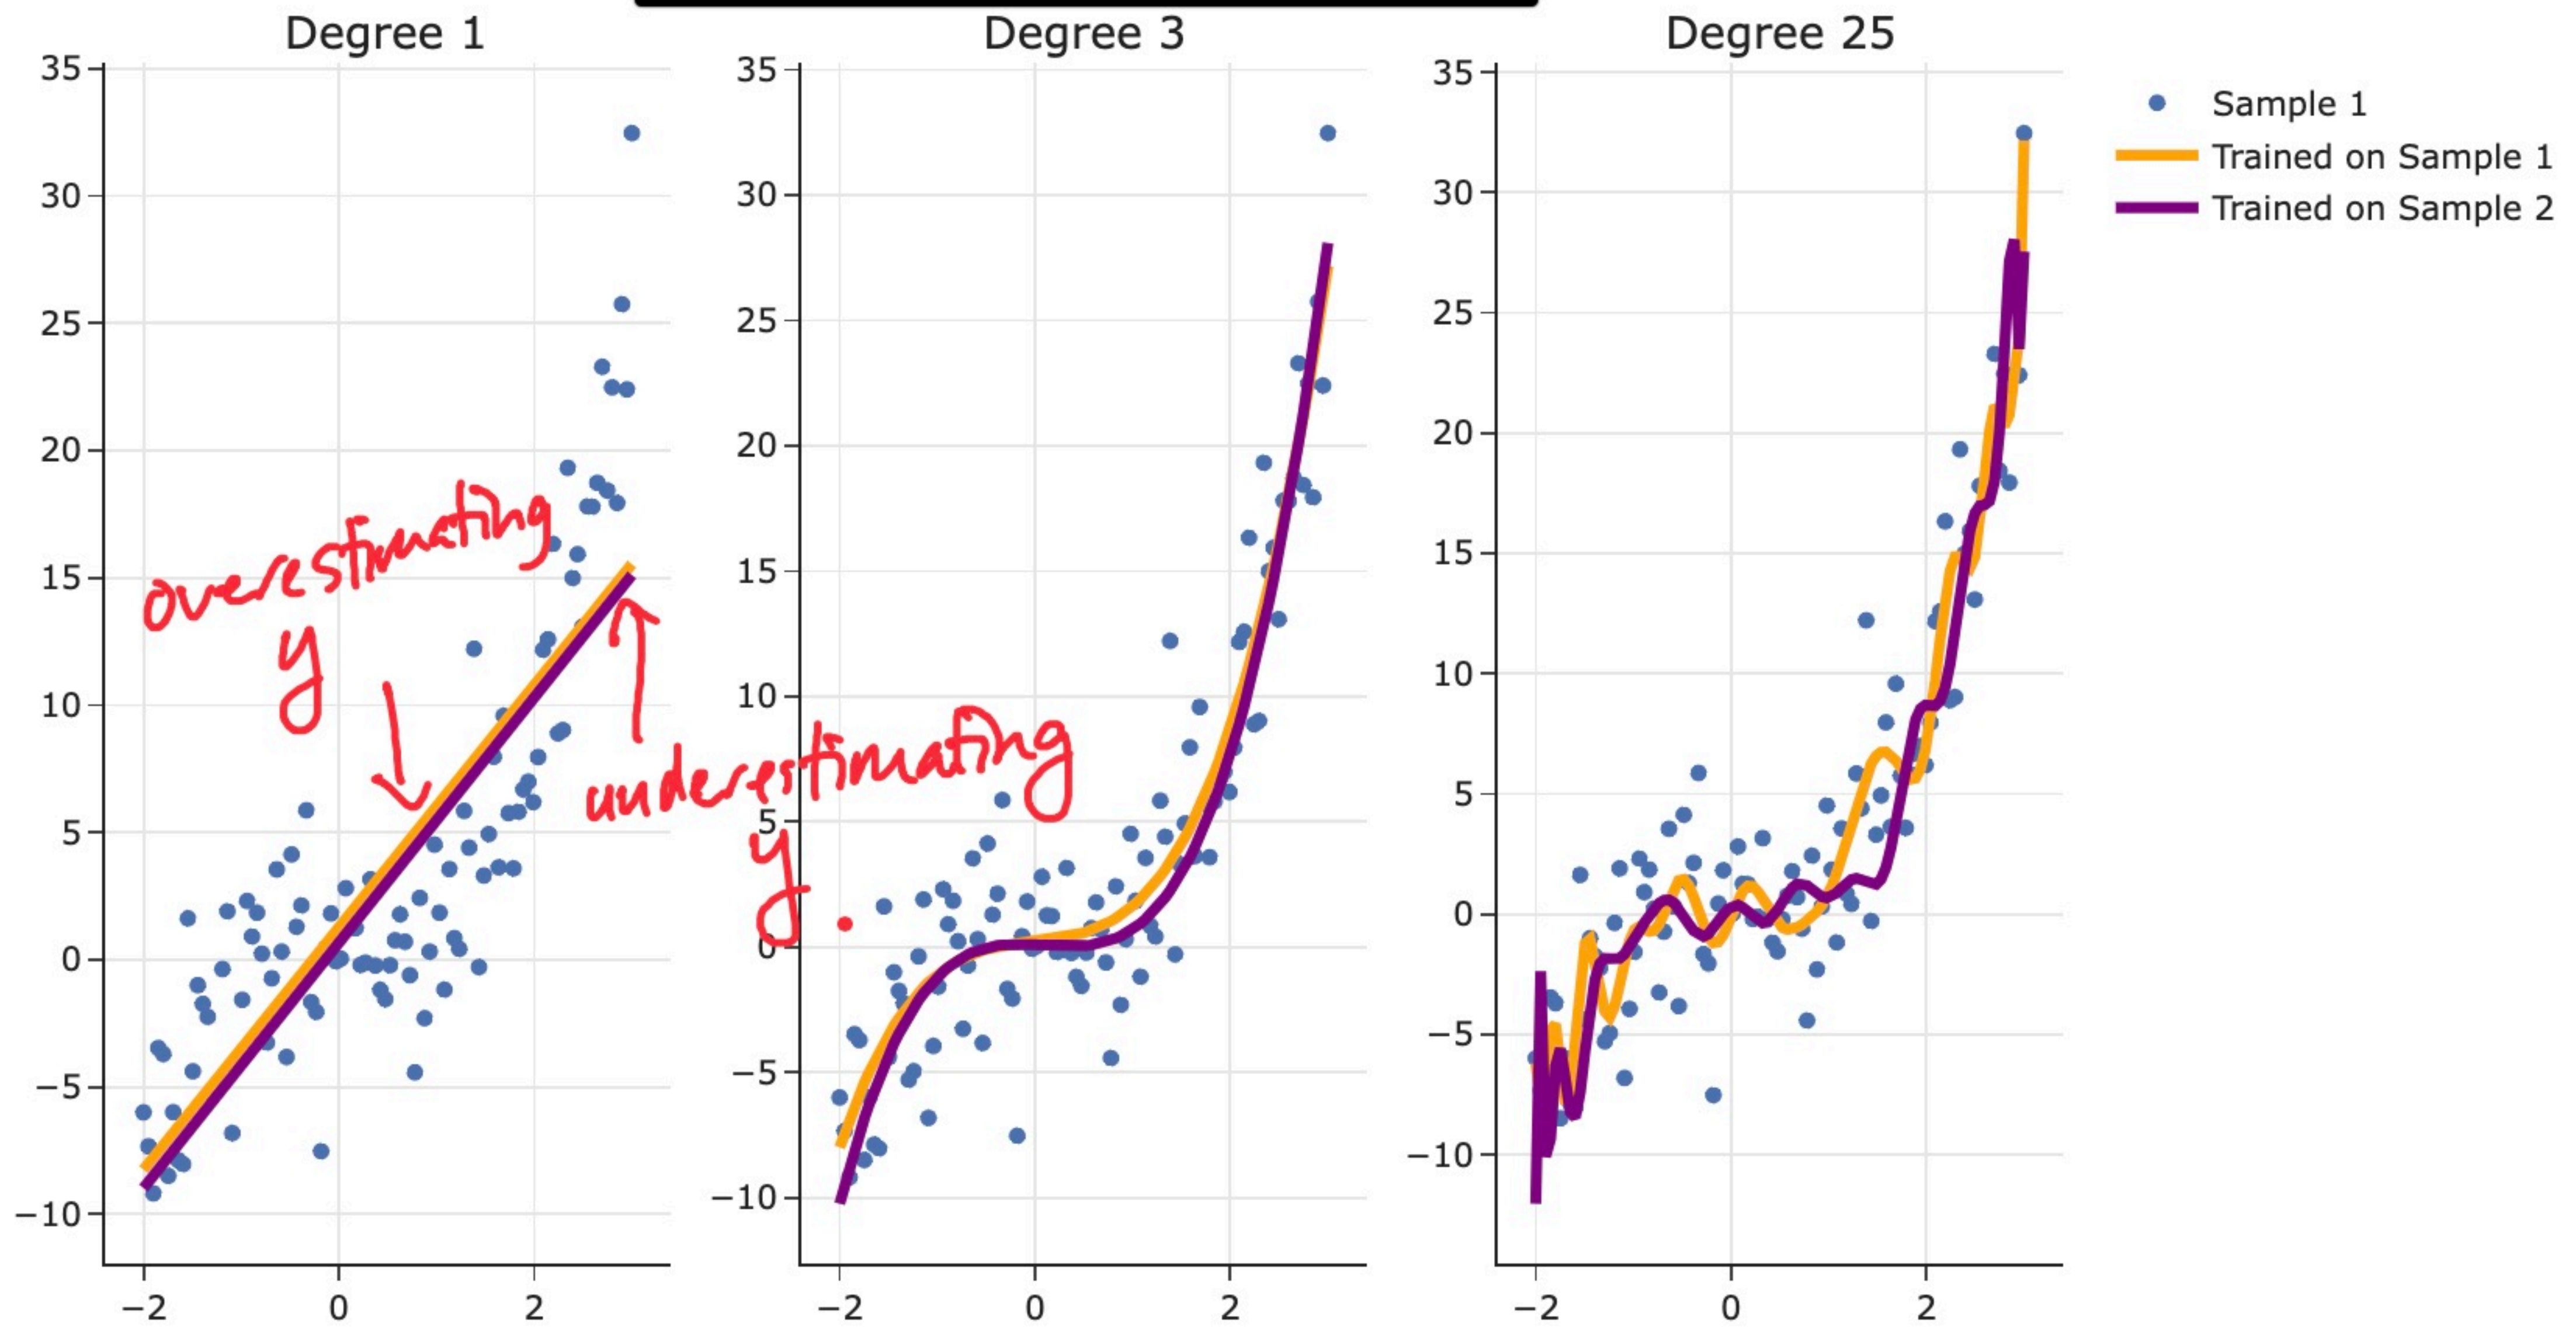


```
In [4]: from sklearn.preprocessing import PolynomialFeatures
# fit_transform fits and transforms the same input.
# We tell it not to add a column of 1s, because
# LinearRegression() does this automatically, later on.
d2 = PolynomialFeatures(3, include_bias=False)
d2.fit_transform(np.array([1, 2, 3, 4, -2]).reshape(-1, 1))
```


```
Out[4]: array([[ 1.,  1.,  1.],
               [ 2.,  4.,  8.],
               [ 3.,  9., 27.],
               [ 4., 16., 64.],
               [-2.,  4., -8.]])
```

x x^2 x^3



• Low bias is good! ✓

ideal



- Here, suppose:
 - The **red bulls-eye** represents your **true weight and height** 🧑.
 - The **dark blue darts** represent **predictions of your weight and height** using different models that were fit using different samples drawn from the same population.

$$\vec{w}^* = \operatorname{argmin}_{\vec{w}} \frac{1}{n} \sum_{i=1}^n (y_i - H(\vec{x}_i))^2$$

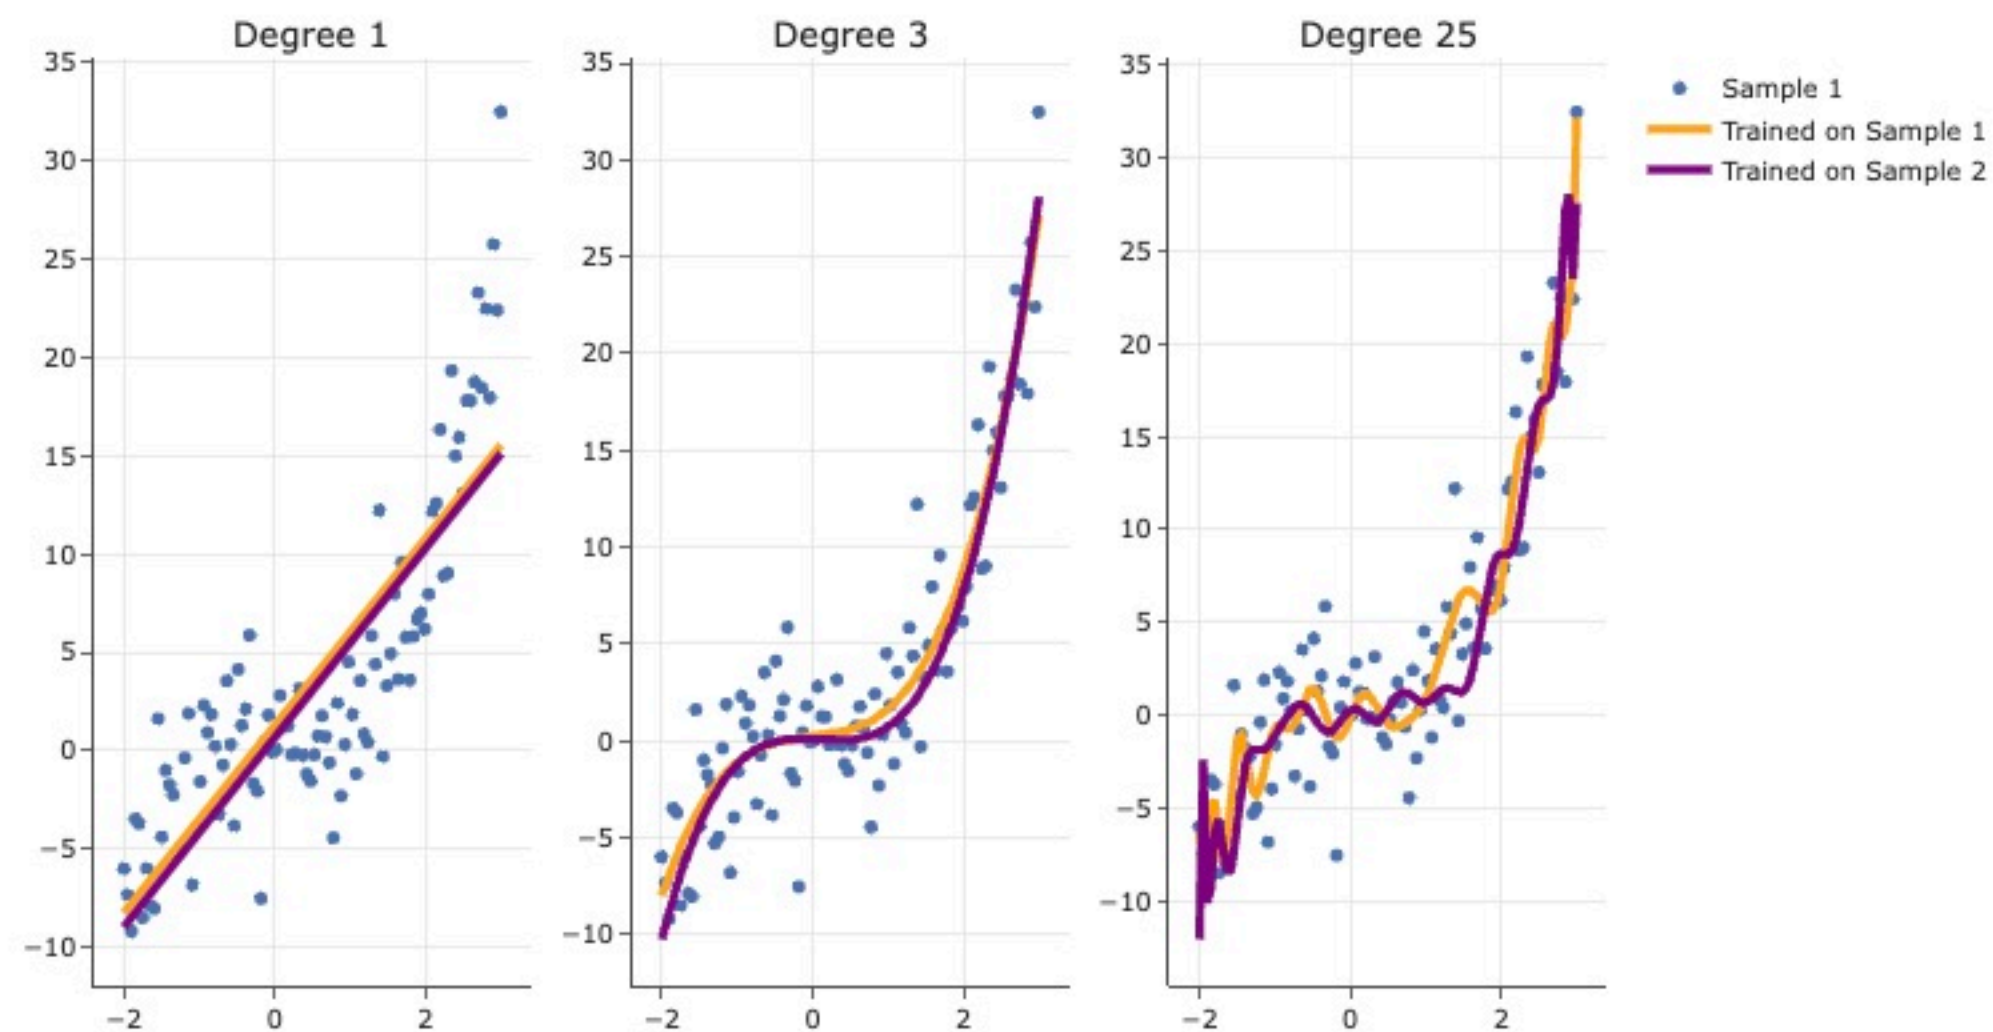
- **Key idea:** A model that works well on past data should work well on future data, if future data looks like past data.
- What we really want is for the:
 - **expected** loss for a new data point $(\vec{x}_{\text{new}}, y_{\text{new}})$,
 - drawn from the same population as the training set, to be small.

That is, we want to minimize **risk**:

$$\text{risk} = \mathbb{E}[y_{\text{new}} - H(\vec{x}_{\text{new}})]^2$$


expected value
random variable

$x_{\text{new}}, y_{\text{new}}$
drawn from same
population
distribution.



- As model variance increases, model bias tends to decrease, and vice versa.

That is, there is a **tradeoff** between bias and variance:

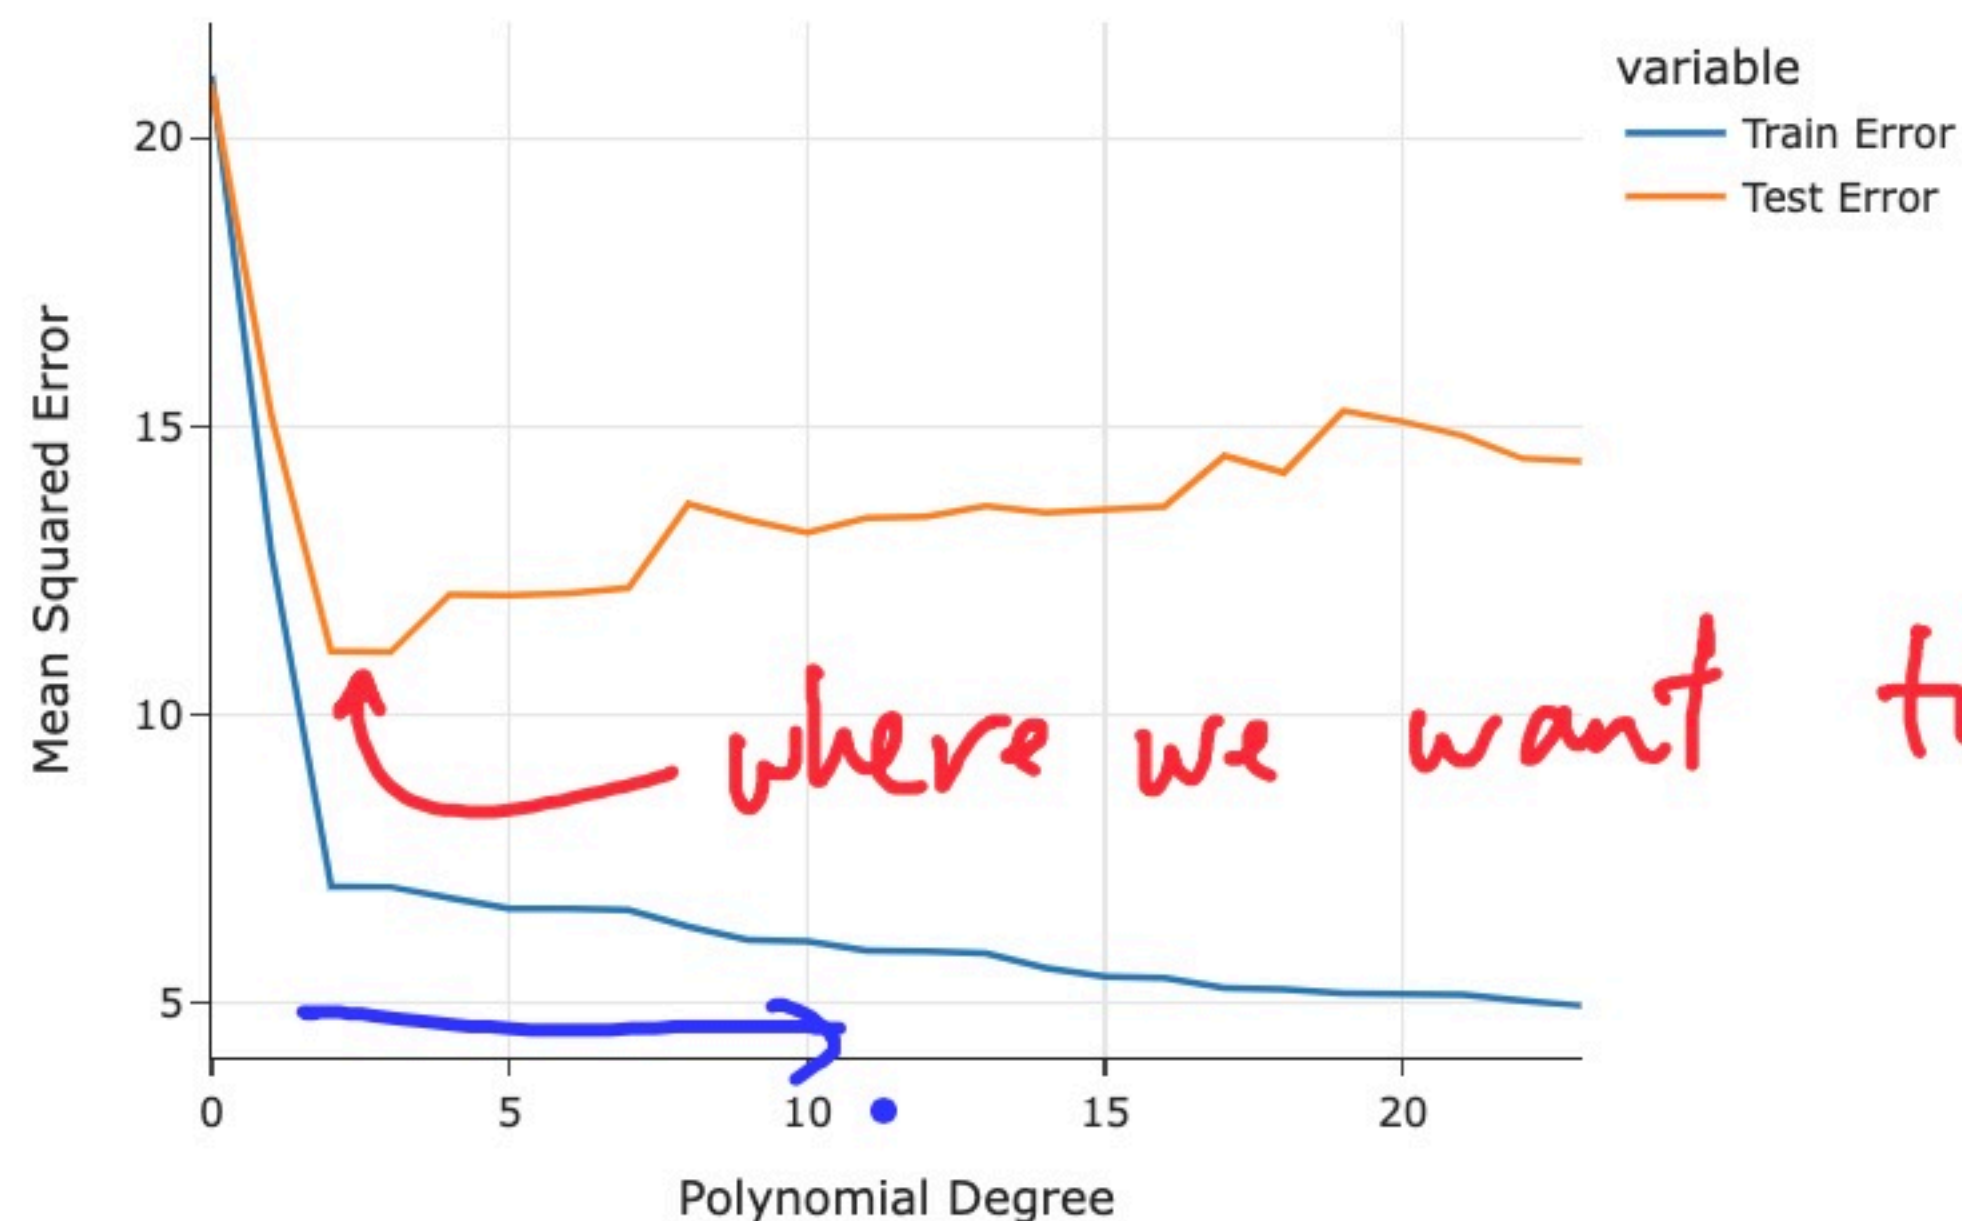


polynomial
degree
=
complexity

- Now that we've performed a train/test split of Sample 1, we'll create models through 25 polynomial features and compute their train and test errors.

```
In [ ]: from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
train_errs = []
test_errs = []
for d in range(1, 26):
    pl = make_pipeline(PolynomialFeatures(d, include_bias=False), LinearRegression())
    pl.fit(X_train, y_train)
    train_errs.append(mean_squared_error(y_train, pl.predict(X_train)))
    test_errs.append(mean_squared_error(y_test, pl.predict(X_test)))
errs = pd.DataFrame({'Train Error': train_errs, 'Test Error': test_errs})
errs
```

```
fig.update_layout(showlegend=True, xaxis_title='Polynomial Degree', yaxis_title='Mean Squared Error')
```




- Training error appears to decrease as polynomial degree increases.
- Test error appears to decrease until a "valley", and then increases again.
- Here, we'd choose a degree of 3, since that degree has the **lowest test error**.

GridSearchCV

- Let's use k -fold cross-validation to choose a polynomial degree that best generalizes to unseen data. As before, we'll choose our polynomial degree from the list $[1, 2, \dots, 25]$.

- GridSearchCV takes in:
 - an **un-fit** instance of an estimator, and
 - a **dictionary** of hyperparameter values to try,

and performs k -fold cross-validation to find the **combination of hyperparameters** with the best average validation performance.



```
In [33]: from sklearn.model_selection import GridSearchCV
GridSearchCV?
```

- Why do you think it's called "grid search"?

Interpreting the results of k -fold cross-validation

- Let's peek under the hood.

```
In [46]: errs_df = util.format_results(searcher)
errs_df
```

Out[46]:

	Degree 1	Degree 2	Degree 3	Degree 4	...	Degree 22	Degree 23	Degree 24	Degree 25
Fold 1	24.36	14.06	7.45	7.62	...	40.48	69.03	43.38	220.49
Fold 2	31.50	19.17	7.95	10.85	...	796.97	1446.11	2208.78	5666.03
Fold 3	15.64	11.78	9.39	9.48	...	61.79	80.64	474.69	4620.71
Fold 4	25.95	17.57	9.47	9.46	...	12.24	12.02	11.86	15.02
Fold 5	11.12	4.74	4.44	4.81	...	8.85	8.54	8.65	8.27


5 rows x 25 columns

validation MSE when using
a degree 3 polynomial
trained on slices 1, 2, 4, 5
(slice 3 = validation!)




```
Out[49]: Degree 3      7.56  
Degree 10     8.33  
Degree 4      8.44  
  
...  
Degree 23    323.27  
Degree 24    549.47  
Degree 25   2106.10  
Length: 25, dtype: float64
```

```
In [50]: fig = errs_df.mean(axis=0).iloc[:18].plot(kind='line', title='Average Validation Error')  
fig.update_layout(xaxis_title='Degree', yaxis_title='Average Validation MSE', showlegend=False)
```



In []: # Chosen automatically by sklearn.


```
.update_layout(xaxis_title='Mean Squared Error', yaxis_title='Model')  
)  
commute_models_summarized
```



smallest average validation MSE.